

Casino Brawyal vs. the Market Analysis: honest status and the road to the proposed model

Written 2026-09-30 for the designer. Sources: docs/Market Analysis.pdf (5 pages, no citations), the codebase (counts verified by reading data/ and src/), the live telemetry service, and web research done today with dates on every number.

Context

The PDF recommends a \$14.99 "ability-array character builder" and lists what a \$14.99 game must contain to earn a Very Positive rating: 3 to 4 classes, 80 to 120 abilities, 30 to 45 minute runs, 15 to 20 enemies across 3 acts, 3 to 6 bosses, 15 to 20 ascension levels, nested tooltips, combat juice, speed toggles, undo, unique icons, controller and Steam Deck support. The designer asked for a blunt comparison of the game as it stands against that model, and a detailed plan to reach it.

Part 1. The PDF itself is not reliable enough to bet the price on

Before comparing the game to the document, the document has to be checked. It cites no sources, and several of its figures are wrong or stale as of today:

PDF says	Verified today	Source
Slay the Spire 2: 91%+ rating, 270,000+ reviews	61% across 200,900 reviews (16 Sep 2026); "reviews turn mixed" coverage	Raijin Steam stats; Kotaku; fixgamingchannel
Luck be a Landlord: 11,000+ reviews	5,331 reviews , 94% positive	Steam store page
Balatro: 110,000+ reviews	198,857 reviews , 98%	Raijin (23 Sep 2026)
Brotato: 75,000+ reviews	118,628 reviews , 96%	Raijin (12 Sep 2026)
Peglin: 82%	84% across 17,616	Raijin (22 Aug 2026)
"nearly 8% of releases reach 1,000 reviews"	5.1% hit rate for roguelike deckbuilders in 2025 ; median game has 360 wishlists , top 10% hold 85% of all wishlists	Venn Studios, 150-game dataset, 2026
"\$10 quality signal filter": under \$10 competes with asset flips	The two closest comparables to this game are under \$10: Luck be a Landlord \$9.99, CloverPit \$9.99, 1,000,000+ copies, a 2-person core team , 90% of 12,734 reviews	Steam, PC Gamer, Panik Arcade
Price band success table (0.5% / 1.4% / 8.8% / 22.7%)	Unsourced. It reads correlation as causation: teams with budget both price higher and succeed more. Higher price does not make a game succeed.	Zukowski's own published finding is that ~500 indie games a year reach 1,000 reviews regardless of release volume

Steam Deck compatibility is "critical"	Card roguelikes: 2.03% of reviews come from Deck-majority players (games over \$1M revenue)	GameDiscoverCo via Game World Observer, 28 Sep 2026
--	--	---

What the PDF gets right: the genre is crowded and hit-driven; the content bar for a premium price is real; readability, juice, speed and undo are the retention drivers; 1,000 reviews is the line where Steam starts promoting you. The scope list on page 5 is a fair description of what \$14.99 buyers expect. Treat that list as the spec. Treat the price as undecided until the content exists and the wishlist count is known.

The relevant cautionary tale is not in the PDF: **Bramble Royale** (GameDiscoverCo, 2025), a well-reviewed roguelike deckbuilder by a developer with a \$1.1M prior hit, launched with 4,117 wishlists and sold **703 copies, \$8,049 gross, in week one**. The author's conclusion: that is the default outcome in this genre unless something drastic lifts you out of the pack. The tag is down year on year.

Part 2. Where the game actually is (verified in the code and telemetry)

PDF expectation for \$14.99	Casino Brawyal today	Gap
3 to 4 classes	1 (Ace). <code>RunState.hero_id</code> is hard-coded; no hero picker; no per-hero ability pool field	2 to 3 heroes, and the plumbing for any second one
80 to 120 unique abilities, 4 to 6 archetypes per class	25 abilities (each with Silver/Gold tiers). Archetypes: Mark/Cash In, Weak synergy, Earn/chips, Block; two or three are viable	55 to 95 abilities; archetype design
30 to 45 minute winning run	Real wins recorded: 13, 19 and 24 minutes (3 of 9 runs). Defeats 4 to 11 minutes	Roughly half the target length. Three acts would fix this on their own
15 to 20 enemies, 3 acts, 3 to 6 bosses	7 enemy defs, 1 act of 10 encounters, 1 boss , 2 elites, 18 lineups, 6 authored paths	Two more acts, 8 to 13 enemies, 2 to 5 bosses
15 to 20 ascension levels	None. <code>hp_mult</code> / <code>dmg_mult</code> exist in the factory, hard-coded to 1.0. No meta-progression, no profile save, no unlocks, no run history	Whole system
Nested colour-coded tooltips	Yes. Keyword panel on cards, status tooltips on units, intent hints	Done
Combat juice	Yes, above the bar. Shake, hitstop, floating numbers, status pips, pose frames, living backdrops, and as of yesterday music and 64 sounds	Pitch escalation on multi-hits is not there (we jitter, not escalate)
1.5x/2x/3x speed toggle	None. ~130 hand-timed waits and tweens	2 to 3 days
Undo before RNG resolves	Partial. A chip can be unsocketed; a fired ability or a turn cannot be undone	2 to 3 days
Distinctive 2D art, unique icons	Yes. 25 of 25 abilities have their own icon; 11 status icons; 7 enemies with pose frames	Done for act 1

Controller, Steam Deck, 14pt fonts	None. No input map, no focus navigation, mouse drag-and-drop everywhere. Card description font floor is 11 px at 1080p, which is ~7 px on a Deck	10 to 14 days plus a font pass
Localization	None. No <code>tr()</code> , ~240 data strings and ~250 UI literals in English	6 to 8 days plus translation
Steam presence	None. No store page, no capsule, no trailer, no demo, no wishlists, no Steamworks	The largest gap of all
Players	4 installs, 16 sessions, 9 runs, 2 hours total since telemetry went live on 21 Sep. All of them are us	Zero external validation

The blunt version. This is a polished vertical slice of one act with one hero and a fifteen-minute run. By the PDF's content bar it is a quarter to a third of a \$14.99 game. Balance has never been tuned against real players: the random bot wins 0 of 40, the four of us win 3 of 9, and the designer chose to ship it hard "to see". It has no tutorial and no onboarding beyond a cinematic. Nobody outside the team has played it, there is no store page accumulating wishlists, and in this genre wishlists take 6 to 12 months to build. If it launched next month at any price the realistic outcome is Bramble Royale's: a few hundred copies.

What is genuinely strong, and it matters because it is what makes the plan below feasible: the architecture. Pure sim, presenters, JSON content with a linter, seeded RNG, a 40-run balance bot, 408 tests, a content pipeline for art and now audio, telemetry with a dashboard. Adding a hero, an act or a heat system is mostly data and a known list of code sites, not a rewrite. The combat feel and the art direction are already at or above the bar the PDF sets. The problem is quantity and market exposure, not quality.

On the hook. CloverPit took the "slot machine roguelike" space a year ago with a horror twist and a million copies. Casino Brawyal's differentiator is the *battler*: a slot machine that pays out chips you socket into a hero's abilities against enemies with intents. That is a real, legible twist on Slay the Spire and it is not what CloverPit or Luck be a Landlord do. It has to be visible in six seconds of capsule and GIF, and today no such asset exists.

Part 3. The plan to reach the proposed model

Ordered by what buys the most risk reduction per week. Engineering days assume one experienced developer with an AI assistant in the repo's TDD rhythm; design and content days are the designer's. Total engineering to the full spec: **roughly 80 to 110 developer-days**, plus content design, art, translation and marketing time. At the current cadence that is 6 to 9 months, which happens to be the minimum wishlist-building window anyway.

Phase 0. Stop building blind (next 3 weeks)

Goal: a Steam page collecting wishlists, twenty external players in telemetry, and the two QoL features every reviewer will mention.

1. **Steam page and Steamworks account.** Capsule art (the machine, Ace, an enemy, chips in flight), six GIFs showing the spin-to-socket-to-strike loop, a 60 second trailer cut from the living backdrop and the barrage. Tagline that names the twist. Tags: Roguelike Deckbuilder, Turn-Based, Card Battler, Casino. Wishlists start counting the day it goes live; every month of delay is a month less at launch.
2. **External playtest, 20 to 30 people**, distributed as the exe we already build, telemetry on (it already records the funnel, lethality per lineup, abilities fired, run length). Read `/api/stats` weekly. This

replaces guessing about the 0/40 curve.

3. **Combat speed toggle** (2 to 3 days). One owner of time scale: `Fx.base_speed`, applied while `combat_screen._busy` is true, restored by `Fx.hitstop` and the ultimate slow-mo instead of the literal 1.0. Fix `HeaderHud` run timer to wall clock. Setting persisted in `AppSettings`.
4. **Undo turn** (2 to 3 days). `CombatSim.snapshot()/restore()` at the start of ASSIGNMENT, covering actors, tray, reel pools, RNG stream states, brain state; an equivalence test over 40 seeds first.
5. **Onboarding**: a first-fight overlay (spin, drag a chip, socket, target, pass) and a codex entry per keyword. Reuse the keyword panel.
6. **Balance pass on act 1 against the playtest data**, not the bot. Target: a first-time player reaches encounter 6 to 8 on average and a competent one wins about a third of runs.

Phase 1. Depth on the hero you have (weeks 4 to 11)

Goal: a demo-ready act 1 with real replayability, for Steam Next Fest (February 2027 or June 2027).

1. **Ace's kit to 45 to 50 abilities**, 5 archetypes (Mark/Cash In, Weak control, Earn engine, Block and counter, Rage and multi-hit). Each new mechanic touches five code sites (`KNOWN_OPS`, `EffectInterpreter`, `AbilityCard._collect_damage`, `GreedyBot._score_effects`, `TelemetryFilter`); budget 3 to 6 new ops. Designer content: 20 to 25 cards.
2. **Profile save** (`user://profile.json`), and **one run-end funnel** `Game.end_run()` replacing the three writers today. Unlocks as data (`data/unlocks/`), a stats screen, seeded run entry, daily run. 7 to 9 days.
3. **Heat 1 to 10** (4 to 6 days): `data/heat/heat_levels.json`, `Heat.modifiers()` folded into `EncounterFactory` (`hp_mult`, `dmg_mult`, gold, rest heal, start HP, `POOL_COPIES`, loan terms). First consolidate the duplicated gold and rest literals between the reward screen and `RunBot`. Win at heat N unlocks N+1 per hero. The bot sweeps heats 0/5/10.
4. **Demo build**: act 1 only, heat off, 3 free ability picks, telemetry on. Demo players wishlist at 2x the rate of page visitors and convert at 18 to 25% at launch.
5. **Creator list and outreach** for the Next Fest window: the roguelike deckbuilder creators who made Balatro and CloverPit (Retromation, Northernlion, Olexa, Wanderbots, DangerouslyFunny) plus the Rogueliker and Turn Based Lovers Next Fest roundups.

Phase 2. Act 2 and a second hero (weeks 12 to 21)

1. **Acts as data** (`data/acts/acts.json`; paths and lineups gain an `act`; `RunState.act_index` and `act_start`; `RunFlow.finish_act` shared by `Game` and `RunBot`; act transition screen; HUD "Act %d %d/%d"; `LivingBackdrop` preset registry; per-act music ids). 4 to 6 days plumbing.
2. **Act 2 content**: 5 to 7 enemies, 1 to 2 bosses, 15 lineups with gold bands pinned in tests, a new backdrop with its UV regions, a combat and boss theme. 10 to 14 days engineering plus art and designer time. Expect 3 to 5 new intent keys or passives, each sim code.
3. **Second hero** from the sheet's Queen tab once the designer finishes it: per-hero ability pool field, hero select screen, `ABILITY_ANIMS` moved into data, `Rewards.upgradeable_pool` filtered by hero, `RunBot.play(hero)`, telemetry `hero` column. 2 to 3 days plumbing plus 5 to 8 days per hero including art from the idle-edit pipeline.

Phase 3. Act 3, third hero, heat to 20, platform (weeks 22 to 34)

1. Act 3 and its boss, third hero, heat 11 to 20: same recipes as phase 2.
2. **Controller and Steam Deck** (10 to 14 days): a pure `PadModel` (tray, cards, enemies, pass; select chip, select card, confirm via `AbilityState.placement_slot`), focus on every screen, prompt glyphs, `pad_*`

input actions, cursor hidden on joypad input. Raise the card description font floor from 11 to at least 15 px, audit every override under 18, add a large-text option. Measure frame time on `gl_compatibility` with the backdrop shaders and offer a reduced-backdrop setting.

3. **Steam plumbing** (7 to 10 days): GodotSteam GDExtension pinned to a 4.5-compatible release, a sixth inert autoload `SteamBridge`, achievements mirrored from the profile, Auto-Cloud for `profile.json` and `run_save.json` (exclude `telemetry/`), Linux export, PCK export filters (today the PCK ships `tests/`, `addons/gut/`, `docs/`, `services/`). Hardware: a Deck or a Linux box is needed; the office network blocks Steam endpoints, so uploads go through the AKS route or another network.
4. **Localization** (6 to 8 days engineering plus translation): string extraction to CSV with stable keys, `Loc` helper so `src/core` never touches `TranslationServer`, keyword numbers templated from `StatusRules`, the "Deal N" test extended per locale, a `tr()` lint in `test_copy_style.gd`, Hebrew and CJK font fallbacks. Which languages: read the `locale` field telemetry already records.
5. **Pitch escalation on multi-hits**: `hit_enemy` pitch stepping per instance within an attack.

Phase 4. Launch decision (month 8 to 9)

Decide Early Access versus 1.0 and the price **from the wishlist count and the content actually shipped**, not from the PDF:

- Under 7,000 wishlists at the decision point: launch at **\$9.99** (the CloverPit and Luck be a Landlord band), Early Access, keep building acts as updates. Fewer sales are needed to reach 1,000 reviews at the lower price only if the game is bought at all; visibility is the constraint, not margin.
- 10,000 or more wishlists and all three acts, three heroes and heat 20 shipped: **\$14.99** is defensible, with the 10 to 15% launch discount the PDF suggests.
- Either way, budget a 2,000 to 5,000 wishlist Next Fest result as the signal that the pitch works.

What the plan does not fix

- **Design bandwidth.** Every item above except the platform work is gated on the sheet: 55 to 95 cards, 8 to 13 enemies, 2 to 5 bosses, 20 heat levels. The code has repeatedly waited on content (the 0.114 to 0.115 gap, the Queen tab still WIP). This is the schedule's critical path.
- **The market does not owe the game an audience.** A Steam page, a demo, creator coverage and a legible hook are the only levers that move the Bramble Royale default. They cost weeks of the designer's time that produce no features, and they are the highest-return work on this list.

Verification

- Every phase ends with `tools/verify.ps1` green, the `RUN BOT STATS` line recorded in `CLAUDE.md`, a commit, a build and a report Artifact, as today.
 - Phase 0 success: Steam page live with a capsule and trailer; 20 external installs visible in `/api/stats`; median first-time run depth and win rate read from the funnel panel.
 - Phase 1 success: demo submitted to a Next Fest; wishlist count and demo-to-wishlist rate recorded.
 - Phase 2 and 3 success: winning runs of 30 to 45 minutes in telemetry; heat 20 clearable by the designer; Deck verification checklist passed on real hardware.
 - Phase 4: the price decision written down with the wishlist number that drove it.
-

Part 4. Who does what: what can be built without the designer

Two questions came with the request for this document: can the missing pieces be added without the designer, and which of them need him. The honest split:

Item	Buildable without the designer	Needs the designer
Combat speed toggle, undo turn	Yes, end to end	Review only
Profile save, stats screen, seeded and daily runs, unlock plumbing	Yes	Which abilities and relics are locked behind what
Heat system plumbing (data shape, factory wiring, unlock rule, bot sweep)	Yes	The 20 levels' actual modifiers and names
Acts plumbing (acts.json, transitions, HUD, ribbon, backdrop registry)	Yes	Act themes, enemy rosters, bosses, backdrop briefs
Hero-select plumbing (per-hero pools, screen, bot, telemetry)	Yes	The heroes: concept, kit, art brief. The Queen tab is WIP
Controller and Steam Deck, font pass, performance	Yes	A Deck or Linux machine to test on
Steam plumbing (GodotSteam, achievements, cloud saves, Linux export)	Yes	The Steamworks account and a network that reaches Steam
Localization engineering (keys, CSV, lint, fonts)	Yes	Which languages; translation budget
Onboarding overlay and codex	Yes, from the existing keyword text	Tone check
Steam page copy, tags, capsule and trailer captures	Drafts, yes	Final voice, capsule art direction, approval
55 to 95 new cards, 8 to 13 enemies, 2 to 5 bosses	Proposals, yes: drafted in the sheet's format with numbers the bot has tested	Every one of them is his call. The doc wins
Balance	Bot runs and telemetry reads	The decision
Price and launch mode	Numbers	The decision
Playtester recruitment, creator outreach	A list and a schedule	The relationships

Proposed working mode, to keep the "doc wins" rule intact: candidate content is drafted into a **Proposals** tab of the capabilities spreadsheet in the sheet's own format, with bot results beside each row. The designer edits or rejects there. Approved rows move to the live tabs and are transcribed with the existing converters. Nothing reaches the shipped JSON that he has not touched.

What starts immediately without waiting on anyone: the Phase 0 engineering (speed, undo, onboarding), the Steam page draft, the trailer and GIF captures, the external playtest build with telemetry, and the profile and heat plumbing from Phase 1. That is four to six weeks of work that needs only review, which is roughly how long the designer needs to write the first batch of cards and the act 2 roster.